

SUSHANT LOKHANDE

San Jose, CA | +1 669-340-6655 | lokhandesushant014@gmail.com | sushantlokhande.me |
linkedin.com/in/sushantlokhande14 | github.com/sushantlokhande14

EXPERIENCE

San Jose State University | Graduate Research Assistant, Applied Systems & Evaluation Aug 2025 to May 2026

- Validated **14 interoperability scenarios** across networking, dependency, and **failure-recovery** paths to confirm recovery behavior before large experiment runs.
- Built an automated evaluation harness for controlled experiments, cutting setup effort **55%** and making every result reproducible from a single config.
- Mentored **200+ students** across **3** courses (Programming Paradigms, AI, ML) through debugging sessions, code reviews, and grading.

Turito | Software Engineering Intern, Reliability & Customer Impact Jun 2022 to Oct 2022

- Kept live classes running for students and teachers across about **12 regions** by cutting session failures **42%** (3.1% to 1.8%) on a **WebRTC** classroom platform.
- Reduced production diagnosis time **35%** by adding structured **metrics**, runtime-state visibility, and targeted logs.
- Cut controlled-disruption failures **42%** (3.1% to 1.8%) by hardening **retries**, timeouts, failure recovery, and distributed state transitions.

Trivia Softwares | Software Engineering Intern Nov 2021 to Jan 2022

- Built a **Rust/Tokio** load-testing framework driving about **25,000** concurrent sessions, exposing lock contention, serialized **WebSocket** writes, and slow session cleanup.
- Cut performance-investigation time about **40%** by building tooling on Rust workers, **Python** aggregation, **SQLite**, Flask, and Chart.js.

TECHNICAL SKILLS

Languages: Python, TypeScript, SQL, C++17, Rust

AI Systems: LLM routing and failover, RAG with source citations and abstention, agent orchestration, human-review routing, evaluation harnesses

Product & Integration: FastAPI, React, REST, gRPC, webhooks, Kafka, Temporal, Docker, AWS

Reliability: failure classification, retries and timeouts, structured telemetry, load testing, root-cause analysis

PROJECTS

Kairo, Voice and Chat Assistant | *React, TypeScript, FastAPI, OpenAI, Vapi, AWS*

- Shipped a chat and voice assistant with LLM routing and call handoff in **under one week**.

Document Action Agent | *React, TypeScript, FastAPI, LLMs*

- Built a document-grounded assistant that verifies sources before recommending actions and routes ambiguous cases to **human review** instead of completing them silently.

Teller, Agent Evaluation Environment | *Browser agents, Claude Code, Python, TypeScript*

- Analyzed **14** computer-use agent runs against an adversarial app (nested frames, missing test IDs, session expiry), classifying each failure as planning, execution, validation, or recovery.

OPEN SOURCE

- NVIDIA/garak:** preserved attempt UUIDs in report aggregation for traceable LLM evals (merged, shipped in v0.17.0)
- corsairdev/corsair:** 6 merged **webhook signature verification** fixes (Telegram, Teams, SharePoint, Monday, GitLab, Jira) that fail closed

EDUCATION

San Jose State University | M.S. in Computer Science | GPA 3.8/4.0 | Spartan Scholar Aug 2024 to May 2026

Coursework: Machine Learning, Artificial Intelligence, Advanced Algorithms, Distributed Computing, Database Systems, GPU Programming

University of Mumbai | B.Tech. in Computer Engineering | GPA 3.9/4.0 Jun 2019 to May 2023

Research: malware image-classification ensemble selected for Dr. Mark Stamp's ML textbook; sole-author contactless palmprint verification (0.9975 AUC)