

SUSHANT LOKHANDE

San Jose, CA | +1 669-340-6655 | lokhandesushant014@gmail.com | sushantlokhande.me |
linkedin.com/in/sushantlokhande14 | github.com/sushantlokhande14

EXPERIENCE

San Jose State University | Graduate Research Assistant, Distributed Systems Aug 2025 to May 2026

- Validated **14 interoperability scenarios** across networking, dependency, and **failure-recovery** paths to confirm recovery behavior before large experiment runs.
- Automated provisioning, configuration, and **telemetry** for containerized test clusters, cutting experiment setup effort **55%** and making every sweep reproducible from config.
- Reduced rebuffering **19%** versus a fixed-rate baseline by building a **Docker**, **WebTransport**, and **Media-over-QUIC** testbed and running controlled adaptation experiments across network and protocol configs.

Turito | Software Engineering Intern, Systems & Reliability Jun 2022 to Oct 2022

- Cut controlled-disruption failures **42%** (3.1% to 1.8%) by hardening **retries**, timeouts, failure recovery, and distributed state transitions.
- Slashed p95 **CPU 22%** by profiling high-frequency media paths and eliminating excess buffer allocation, payload copies, repeated routing lookups, and **lock contention**.
- Reduced production diagnosis time **35%** by adding structured **metrics**, runtime-state visibility, and targeted logs.

Trivia Softwares | Software Engineering Intern Nov 2021 to Jan 2022

- Built a **Rust/Tokio** load-testing framework driving about **25,000** concurrent sessions, exposing lock contention, serialized **WebSocket** writes, and slow session cleanup.
- Raised equivalent-load signaling throughput about **31%** by surfacing concurrency bottlenecks that drove targeted backend fixes.
- Cut performance-investigation time about **40%** by building tooling on Rust workers, **Python** aggregation, **SQLite**, Flask, and Chart.js.

TECHNICAL SKILLS

Languages: Python, C++17, Rust, Java, TypeScript, SQL

Backend & Distributed: FastAPI, gRPC, Kafka, Temporal, REST, asyncio, concurrency, caching, rate limiting, authorization gateways

Infrastructure: Docker, Kubernetes, AWS, Linux, CI/CD, Git, profiling and load testing

Data: PostgreSQL, MongoDB, Neo4j, SQLite, Qdrant, Pandas, NumPy

PROJECTS

Portcullis, Policy Gateway | Rust

- Built a **Rust** proxy that centralizes authorization and rate limiting and emits auditable accept, reject, and throttle decisions.

Relay, LLM Inference Gateway | Python, FastAPI, asyncio, SQLite, SSE

- Cut p50 latency from **759 ms to 44 ms** over a **4,000**-request load test with an OpenAI-compatible gateway doing exact and **semantic caching**, failover, and rate limiting.

Proxima, Vector Search Engine | C++17, AVX2/SIMD, HNSW, pybind11

- Reached **0.999 recall** at **12,049 QPS** single-threaded, **1.8x** hnswlib and **2.5x** FAISS, with an HNSW engine written from scratch in C++17.

OPEN SOURCE

- corsairdev/corsair**: 6 merged **webhook signature verification** fixes (Telegram, Teams, SharePoint, Monday, GitLab, Jira) that fail closed
- fission/fission**: 3 merged PRs, Kubernetes event and KEDA object capture in support dumps, removed duplicate package enumeration

EDUCATION

San Jose State University | M.S. in Computer Science | GPA 3.8/4.0 | Spartan Scholar Aug 2024 to May 2026

Coursework: Distributed Computing, Computer Networks, Database Systems, Advanced Algorithms, Machine Learning, Artificial Intelligence

University of Mumbai | B.Tech. in Computer Engineering | GPA 3.9/4.0 Jun 2019 to May 2023

Research: malware image-classification ensemble selected for Dr. Mark Stamp's ML textbook; sole-author contactless palmprint verification (0.9975 AUC)