

SUSHANT LOKHANDE

San Jose, CA | +1 669-340-6655 | lokhandesushant014@gmail.com | sushantlokhande.me |
linkedin.com/in/sushantlokhande14 | github.com/sushantlokhande14

EXPERIENCE

San Jose State University | Graduate Research Assistant, Distributed Systems & Data Aug 2025 to May 2026

- Built an automated **telemetry pipeline** that collected, normalized, and analyzed experiment data end to end, cutting setup and analysis effort **55%**.
- Validated **14 interoperability scenarios** across networking, dependency, and **failure-recovery** paths to confirm recovery behavior before large experiment runs.
- Proved a **2.4 pt** ensemble gain was real with **McNemar** testing at $p < 0.001$ instead of reporting a single accuracy number, across **1,700** held-out samples.

Turito | Software Engineering Intern, Distributed Systems Jun 2022 to Oct 2022

- Slashed p95 **CPU 22%** by profiling high-frequency media paths and eliminating excess buffer allocation, payload copies, repeated routing lookups, and **lock contention**.
- Reduced production diagnosis time **35%** by adding structured **metrics**, runtime-state visibility, and targeted logs.
- Lowered network egress **14%** and p95 CPU **22%** on a high-frequency event path by removing redundant processing, payload copies, and per-packet logging overhead.

Trivia Softwares | Software Engineering Intern Nov 2021 to Jan 2022

- Cut investigation time about **40%** by building a pipeline that aggregated load-test output from Rust workers into **SQLite** with Python and served comparisons in Flask.
- Built a **Rust/Tokio** load-testing framework driving about **25,000** concurrent sessions, exposing lock contention, serialized **WebSocket** writes, and slow session cleanup.

TECHNICAL SKILLS

Languages: Python, SQL, C++17, Rust, Java

Data Systems: Hadoop MapReduce, Hive, PostgreSQL, MongoDB, Neo4j, SQLite, Qdrant, Pandas, NumPy

Pipelines: Kafka, Temporal, FastAPI, telemetry collection and normalization, reproducible experiment configs, load testing

Infrastructure: Docker, Kubernetes, AWS, Linux, CI/CD, Git, profiling

PROJECTS

Melody Metrics, Big-Data Pipeline | C++, Hadoop, MapReduce, Hive, Python, XGBoost

- Processed **1M+** heterogeneous records through C++ **MapReduce** and **Hive** into a curated **140,000**-row dataset powering an **84%**-accurate 81-genre classifier.

Graph Connect, Graph-Backed Platform | FastAPI, Neo4j, MongoDB, React, Docker, AWS

- Modeled **50,000+** relationships in **Neo4j** with two-hop Cypher recommendations and load-tested APIs with Locust at about **1,200** users against sub-300 ms targets.

Proxima, Vector Search Engine | C++17, AVX2/SIMD, HNSW, pybind11

- Reached **0.999 recall** at **12,049 QPS** single-threaded, **1.8x** hnswlib and **2.5x** FAISS, with an HNSW engine written from scratch in C++17.

OPEN SOURCE

- **fission/fission**: 3 merged PRs, Kubernetes event and KEDA object capture in support dumps, removed duplicate package enumeration
- **google/osv-scalibr**: upstream change to Google's software composition analysis library (merged)

EDUCATION

San Jose State University | M.S. in Computer Science | GPA 3.8/4.0 | Spartan Scholar Aug 2024 to May 2026

Coursework: Database Systems, NoSQL, Distributed Computing, Advanced Algorithms, Machine Learning, Computer Networks

University of Mumbai | B.Tech. in Computer Engineering | GPA 3.9/4.0 Jun 2019 to May 2023

Research: malware image-classification ensemble selected for Dr. Mark Stamp's ML textbook; sole-author contactless palmprint verification (0.9975 AUC)