

SUSHANT LOKHANDE

San Jose, CA | +1 669-340-6655 | lokhandesushant014@gmail.com | sushantlokhande.me |
linkedin.com/in/sushantlokhande14 | github.com/sushantlokhande14

EXPERIENCE

San Jose State University | Graduate Research Assistant and Teaching Assistant Aug 2025 to May 2026

- Mentored **200+ students** across **3** courses (Programming Paradigms, AI, ML) through debugging sessions, code reviews, and grading.
- Built an automated evaluation harness for controlled experiments, cutting setup effort **55%** and making every result reproducible from a single config.
- Reduced rebuffering **19%** versus a fixed-rate baseline by building a **Docker**, **WebTransport**, and **Media-over-QUIC** testbed and running controlled adaptation experiments across network and protocol configs.
- Boosted malware classification to **80.2%** across 17 families with a soft-voting **PyTorch** ensemble (CNN, ResNet50, ViT), **2.4 pts** over the best single model at $p < 0.001$.

Turito | Software Engineering Intern, EdTech Platform Jun 2022 to Oct 2022

- Kept live classes running for students and teachers across about **12 regions** by cutting session failures **42%** (3.1% to 1.8%) on a **WebRTC** classroom platform.
- Slashed p95 **CPU 22%** by profiling high-frequency media paths and eliminating excess buffer allocation, payload copies, repeated routing lookups, and **lock contention**.
- Reduced production diagnosis time **35%** by adding structured **metrics**, runtime-state visibility, and targeted logs.
- Lowered network egress **14%** by removing redundant media-path processing and simplifying routing and payload handling.

Trivia Softwares | Software Engineering Intern Nov 2021 to Jan 2022

- Cut performance-investigation time about **40%** by building tooling on Rust workers, **Python** aggregation, **SQLite**, Flask, and Chart.js.
- Built a **Rust/Tokio** load-testing framework driving about **25,000** concurrent sessions, exposing lock contention, serialized **WebSocket** writes, and slow session cleanup.

TECHNICAL SKILLS

AI-Native Workflow: Claude Code, OpenAI Codex, GitHub Copilot, agent evaluation, test-first acceptance criteria, diff review

Languages: Java, TypeScript, JavaScript, Python, SQL, C++

Stack: React, Next.js, FastAPI, REST, PostgreSQL, SQLite, Redis, Docker, AWS

Testing: pytest, Selenium, Locust, CI/CD, Git

PROJECTS

Teller, Agent Evaluation Environment | *Browser agents, Claude Code, Python, TypeScript*

- Analyzed **14** computer-use agent runs against an adversarial app (nested frames, missing test IDs, session expiry), classifying each failure as planning, execution, validation, or recovery.
- Ported the same workflows to a second tenant with a **16-line** overlay and zero re-recording by designing a reusable capability layer.

AI Learning Hub, Grounded RAG | *Python, FastAPI, LangChain, OpenAI, Docker, AWS*

- Grounded answers over **10,000+** articles with source citations and abstention when retrieved evidence was insufficient.

AutoGrade AI, Multi-Agent Evaluation Platform | *Python, Temporal, Kafka, gRPC, PostgreSQL, Qdrant*

- Orchestrated **4-dimension** rubric grading (correctness, efficiency, resilience, quality) across multiple LLM agents on **Temporal** and **Kafka**, with confidence gates and human-review queues.

OPEN SOURCE

- **NVIDIA/garak**: preserved attempt UUIDs in report aggregation for traceable LLM evals (merged, shipped in v0.17.0)
- **microsoft/markitdown**: fixed EPUB conversion for percent-encoded manifest paths, with regression tests (merged)

EDUCATION

San Jose State University | M.S. in Computer Science | GPA 3.8/4.0 | Spartan Scholar Aug 2024 to May 2026

Coursework: Artificial Intelligence, Machine Learning, Database Systems, Distributed Computing, Advanced Algorithms, NoSQL

University of Mumbai | B.Tech. in Computer Engineering | GPA 3.9/4.0 Jun 2019 to May 2023

Research: malware image-classification ensemble selected for Dr. Mark Stamp's ML textbook; sole-author contactless palmprint verification (0.9975 AUC)