

SUSHANT LOKHANDE

San Jose, CA | +1 669-340-6655 | lokhandesushant014@gmail.com | sushantlokhande.me | linkedin.com/in/sushantlokhande14 | github.com/sushantlokhande14

EXPERIENCE

San Jose State University | Graduate Research Assistant, ML Systems Aug 2025 to May 2026

- Automated provisioning, configuration, and **telemetry** for containerized test clusters, cutting experiment setup effort **55%** and making every sweep reproducible from config.
- Boosted malware classification to **80.2%** across 17 families with a soft-voting **PyTorch** ensemble (CNN, ResNet50, ViT), **2.4 pts** over the best single model at $p < 0.001$.
- Reduced rebuffering **19%** versus a fixed-rate baseline by building a **Docker**, **WebTransport**, and **Media-over-QUIC** testbed and running controlled adaptation experiments across network and protocol configs.
- Generated **136,000** training images from **17,000** binaries across **8** visual representations and benchmarked **6** model families on **1,700** held-out samples.

Turito | Software Engineering Intern, Platform Performance Jun 2022 to Oct 2022

- Slashed p95 **CPU 22%** by profiling high-frequency media paths and eliminating excess buffer allocation, payload copies, repeated routing lookups, and **lock contention**.
- Cut controlled-disruption failures **42%** (3.1% to 1.8%) by hardening **retries**, timeouts, failure recovery, and distributed state transitions.
- Reduced production diagnosis time **35%** by adding structured **metrics**, runtime-state visibility, and targeted logs.
- Prototyped a **Rust** + **QUIC** media path serving about **8,000** lightweight subscribers with about **24%** lower startup time.

Trivia Softwares | Software Engineering Intern Nov 2021 to Jan 2022

- Built a **Rust/Tokio** load-testing framework driving about **25,000** concurrent sessions, exposing lock contention, serialized **WebSocket** writes, and slow session cleanup.
- Raised equivalent-load signaling throughput about **31%** by surfacing concurrency bottlenecks that drove targeted backend fixes.

TECHNICAL SKILLS

Languages: Python, C++17, Rust, TypeScript, Java, SQL

AI/ML Infrastructure: LLM gateways, semantic caching, provider failover, HNSW vector search, PyTorch, Hugging Face, sentence-transformers

Backend & Platform: FastAPI, asyncio, gRPC, Kafka, Temporal, Redis, PostgreSQL, SQLite, Docker, AWS, SSE

Search & Ranking: two-tower retrieval, ANN candidate generation, neural ranking, Recall/NDCG/MAP evaluation

PROJECTS

Relay, LLM Inference Gateway | *Python, FastAPI, asyncio, SQLite, SSE*

- Cut p50 latency from **759 ms to 44 ms** over a **4,000**-request load test with an OpenAI-compatible gateway doing exact and **semantic caching**, failover, and rate limiting.
- Served **78%** of requests from cache and cut modeled provider cost **76%**, with live p50/p95, token, and cost metrics.

Proxima, Vector Search Engine | *C++17, AVX2/SIMD, HNSW, pybind11*

- Reached **0.999 recall** at **12,049 QPS** single-threaded, **1.8x** hnswlib and **2.5x** FAISS, with an HNSW engine written from scratch in C++17.

ReelRank, Two-Stage Recommender | *PyTorch, C++ HNSW, FastAPI, React, TypeScript*

- Built a **two-tower** retrieval model with ANN candidate generation and a neural ranker, evaluated on a leakage-free temporal split with Recall, NDCG, and MAP.

OPEN SOURCE

- NVIDIA/garak**: preserved attempt UUIDs in report aggregation for traceable LLM evals (merged, shipped in v0.17.0)
- fission/fission**: 3 merged PRs, Kubernetes event and KEDA object capture in support dumps, removed duplicate package enumeration

EDUCATION

San Jose State University | M.S. in Computer Science | GPA 3.8/4.0 | Spartan Scholar Aug 2024 to May 2026

Coursework: Machine Learning, Distributed Computing, Database Systems, Advanced Algorithms, Artificial Intelligence, NoSQL

University of Mumbai | B.Tech. in Computer Engineering | GPA 3.9/4.0 Jun 2019 to May 2023

Research: malware image-classification ensemble selected for Dr. Mark Stamp's ML textbook; sole-author contactless palmprint verification (0.9975 AUC)