

SUSHANT LOKHANDE

San Jose, CA | +1 669-340-6655 | lokhandesushant014@gmail.com | sushantlokhande.me |
linkedin.com/in/sushantlokhande14 | github.com/sushantlokhande14

EXPERIENCE

San Jose State University | Graduate Research Assistant, ML & Distributed Systems Aug 2025 to May 2026

- Built and evaluated **RAG** pipelines (chunking, embeddings, vector retrieval, top-k, grounded prompts), separating **retrieval** failures from generation failures to guide fixes.
- Reduced rebuffering by **19%** versus a fixed-rate baseline by building a **Docker**, **WebTransport**, and **Media-over-QUIC** testbed and running controlled adaptation experiments across network and protocol configurations.
- Cut experiment setup effort by about **55%** by automating orchestration, configuration, **telemetry** collection, and analysis, making large experiment sweeps reproducible.
- Supported **200+ students** as a TA in Programming Paradigms, AI, and Machine Learning through debugging sessions, code review, and grading.

Turito | Software Engineering Intern, Real-Time Product Jun 2022 to Oct 2022

- Cut controlled-disruption failures by **42%** (3.1% to 1.8%) by hardening **retries**, timeouts, failure recovery, and distributed state transitions.
- Slashed p95 **CPU** by **22%** by profiling high-frequency media paths and removing excess buffer allocation, payload copies, repeated routing lookups, and **lock contention**.
- Reduced production diagnosis time by **35%** by adding structured **metrics**, runtime-state visibility, and targeted logs.
- Lowered network egress by **14%** by removing redundant media-path processing and simplifying routing and payload handling.

Trivia Softwares | Software Engineering Intern Nov 2021 to Jan 2022

- Built a **Rust/Tokio** load-testing framework driving about **25,000** concurrent sessions, exposing lock contention, serialized **WebSocket** writes, and slow session cleanup.
- Reduced performance-investigation time by about **40%** with tooling built on Rust workers, **Python** aggregation, **SQLite**, Flask, and Chart.js.

TECHNICAL SKILLS

Languages: Python, TypeScript, JavaScript, C++, SQL

LLM & AI: RAG, embeddings, evals, agents, tool use, LangChain, OpenAI and Anthropic APIs, PyTorch, Hugging Face

Backend & Frontend: FastAPI, asyncio, SSE, REST, gRPC, React, Next.js, Node.js

Data & Infra: PostgreSQL, Redis, Qdrant, SQLite, Kafka, Temporal, Docker, AWS, pytest

PROJECTS

Relay, LLM Inference Gateway | *Python, FastAPI, asyncio, SQLite, SSE*

- Built an OpenAI-compatible gateway with exact and **semantic caching**, provider **failover**, and rate limiting that cut p50 latency from **759 ms to 44 ms** over a 4,000-request load test.

ReelRank, Two-Stage Recommender | *PyTorch, C++ HNSW, FastAPI, React, TypeScript*

- Built a two-tower retrieval model with ANN candidate generation and a neural ranker, served through **FastAPI** and a **React/TypeScript** UI with metadata-aware cold start.

Teller, Agent Evaluation Environment | *Browser agents, Python, TypeScript*

- Built an adversarial target app (nested frames, missing test IDs, session expiry, permission denials) and analyzed **14 agent runs**, classifying failures as planning, execution, validation, or recovery.

OPEN SOURCE

- **microsoft/markitdown**: fixed EPUB conversion for percent-encoded manifest paths, with regression tests (merged)
- **NVIDIA/garak**: preserved attempt UUIDs in report aggregation for traceable LLM evals (merged, shipped in v0.17.0)

EDUCATION

San Jose State University | M.S. in Computer Science | GPA 3.8/4.0 | Spartan Scholar Aug 2024 to May 2026

Coursework: Machine Learning, Artificial Intelligence, Distributed Computing, Database Systems, Advanced Algorithms, NoSQL

University of Mumbai | B.Tech. in Computer Engineering | GPA 3.9/4.0 Jun 2019 to May 2023

Research: malware image-classification ensemble selected for Dr. Mark Stamp's ML textbook; sole-author contactless palmprint verification (0.9975 AUC)